

Distributed Data Store

Large-Scale Distributed file system

- Q: What if we have too much data to store in a single machine?
- Q: How can we create one big filesystem over a cluster of machines, whose data is reliable despite individual node failures?
- Basic idea:
 - Split the data into small chunks
 - Distribute the chunks to nodes
 - Replicate each chunk at multiple nodes to deal with failure
 - A master server with the filename -> chunk mapping
- GFS (HDFS)

```
[ GFS diagram ]  
  
  client  
  
  master  
  
chunkserver chunkserver chunkserver
```

- Provides a file system (global namespace) over cluster of machines
- A file is split into (often) 64M chunks
 - * each chunk is replicated on multiple chunkservers
- GFS master + chunk server
- Master provides
 - * file name -> chunk server mapping
 - * chunk server monitoring and replacement
- Chunk server provides actual data chunk

- * best chunkserver is chosen to minimize network bandwidth consumption
- Optimized for
 - * Sequential/random read of large file
 - * Append at the end
 - Synchronizing random writes much more costly with failures
- Q: Is a single master server enough to deal with the huge filesystem?
 - * Ex) 1 billion files. 64 byte name per file. 4GB per file. How much space to remember the mapping?
- SAN (Storage Area Network)
 - “Scale up” approach for disks
 - Provides SCSI interface over many hard disks (often thru optical LAN)
 - All nodes in the cluster see one shared big disk
 - Convenient but very expensive

NoSQL datastores

- Limitation of relational databases
 - Relation: very general. But is it the best data representation?
 - * e.g., java objects
 - SQL: very expressive query language. But is this expressiveness necessary?
 - * e.g., student records. How will it be accessed?
 - The power of SQL also leads to complexity. Do we really need it?
 - * Very difficult to reimplement full relational database on a cluster architecture
 - * Need for a new system that is designed from the ground up for
 - Ultimate scalability for highly distributed environment
 - Core functionality necessary for new Web apps
 - ACID: very strong data integrity guarantee. But is it necessary?
 - * e.g., user’s status updates. Is ACID necessary?
- Difficulty of maintaining a shared state among distributed nodes
 - **Byzantine-general problem**
 - * Two generals A, B need to attack together to win

- * Messenger may be lost during delivery (say with 90% chance)
- * Q: How can they make sure simultaneous attack on Sunday midnight?
- * Q: Any problem?
 - A sends a messenger to B
 - A attacks after sending a messenger
 - B attacks once he hears from a messenger
- * Q: Any better?
 - A sends a messenger to B
 - B sends ACK messenger back to A
 - A attacks if he gets ACK from B
 - B attacks if he hears from A's messenger
- * Q: Any better?
 - A sends a messenger. B ACKS. A ACKS back to B.
- * Q: How to improve chance of success?
- Remarks
 - * Keeping states on multiple nodes synchronized is often VERY COSTLY.
 - Synchronization overhead outweighs benefit for more than a few nodes
 - * If possible, relax “state synchronization guarantee” on multiple nodes
 - Otherwise, the overall performance will be too based compared to centralized processing
- CAP theorem (by Eric Brewer)
 - * *Consistency*: after an update, all readers will see the latest update
 - * *Availability*: continuous operation despite node failure
 - * *Partition Tolerance*: continuous operation despite network partition
 - * Only two of three characteristics can be guaranteed by any system
- Two important questions:
 1. Q: What data model to use?
 - *(Key, value) store*
 - *Column store*
 - *Document store*
 2. Q: How to relax consistency guarantee?
 - MongoDB
 - * Atomicity guarantee on individual operation

- * No transaction support on multiple operations
 - * The application is responsible for concurrency control
- ACID vs. BASE
 - * *Basically Available*: system works basically all the time
 - * *Soft-state*: does not have to be consistent all the time
 - * *Eventual consistency*: will be in a consistent state eventually
- Possible consistency models
 - * Read Your Own Writes (RYOW) Consistency
 - * Session Consistency
 - * Monotonic-Read Consistency: readers will see only newer update in the future
 - * ...
- **(Key, value) store**
 - Example: Amazon SimpleDB, Redis, Memcached, ...
 - Example scenario
 - * Key: student id, Value: (student name, address, email, major, ...)
 - * Q: How to process “find all student info with sid 301”?

 - * Q: How to process “find student email with sid 301”?

 - * Q: How to process “change the major of sid 301 to CS”?

 - * Q: How to process “find student whose email is cho@cs”?

 - * Remarks
 - Read can be inefficient for a single-field retrieval
 - Update can be inefficient
 - Data access using a non-primary key is not supported

- Need to create and maintain a separate index for non-key access

- **Column store**

- Example: Google's BigTable, Cassandra (Facebook), ...
- Each data item is a row in a table - Each row has a unique row-key, but may not have all column values
- Columns are grouped as "column families"
 - * Preserves locality of column access and storage
- *Data is accessed through (row-key, column-name) value*
- Compared to (key,value) store, more structure of the data is exposed to the system
 - * More efficient update and retrieval
- Q: What about retrieval based on column values only without row-key (like student by email)?
 - * Data access is still limited to row-key

- **Document store**

- Example: MongoDB, Amazon DynamoDB, CouchDB
- Document consists of multiple fields, which are (key, value) pairs.
- No preset "schema" for documents
- Example document:

```
Title: CouchDB
Categories: [Database, NoSQL, Document Database]
Date: Jan 10, 2011
Author: ...
```

- User specifies the fields on which to build an index
 - * Allows data retrieval based on "non-key" fields (by using appropriate indexes)
 - * But, this data store will be more complex to build than others and less scalable

Consistent hashing

- Q: How to distribute objects with different keys to k nodes? What about “ $\text{hash}(\text{object}) \bmod k$ ”?
- Q: How to minimize data reorganization when a new node is introduced?
 - *Consistent hashing*
 - * Hash BOTH objects and *nodes* to a hash ring
 - * Assign objects to the node right behind them
 - Q: What to do when a new node is introduced?
- Q: Non-uniform hash range for each node. How to minimize load imbalance?
 - Hash many virtual nodes to the ring
 - Assign virtual nodes to physical nodes
- Q: How to make data available despite node failure?
 - Replication to the next k nodes

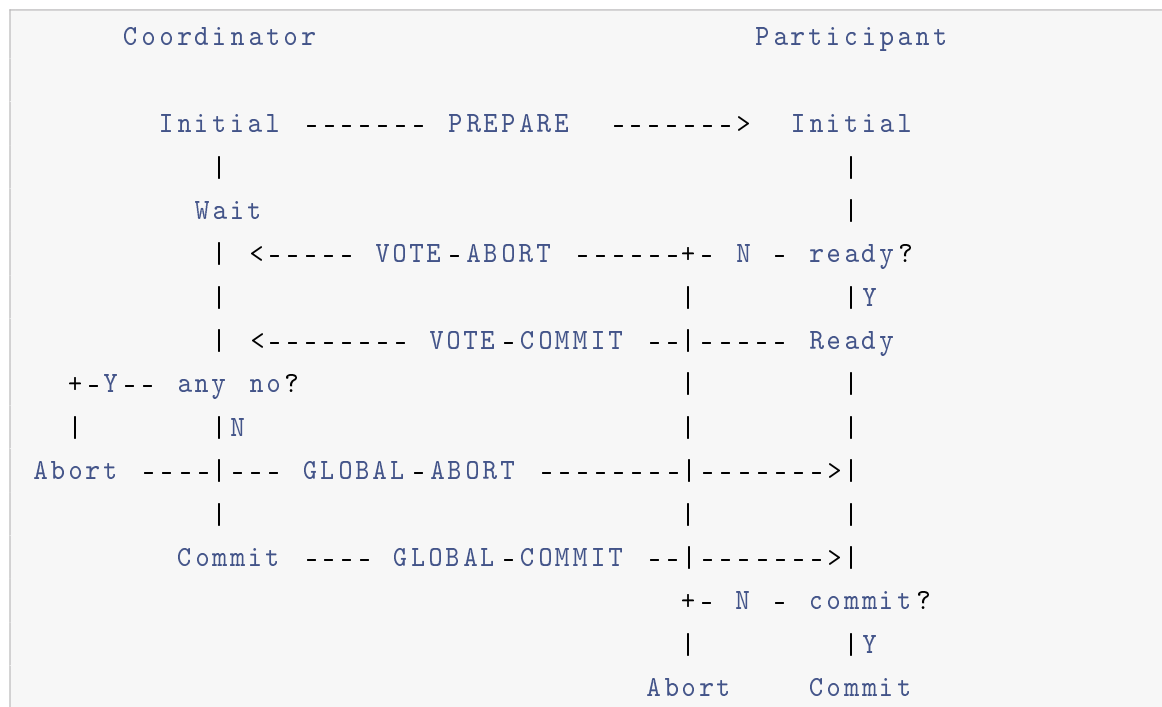
Distributed Transactions

- Q: How can we guarantee atomicity of a transaction on multiple machines? Either everyone commits or no one should commit.

Two-phase commit

- Before commit, ask everyone whether they are ready
 - PREPARE -> VOTE-COMMIT/VOTE-ABORT

- Note: Anyone who said ready cannot say otherwise later
- If everyone says yes, commit
 - if anyone says no, abort



- Q: Any potential problem of two phase commit?
- Q: Where should we add timeout to avoid indefinite wait?
 - Remark: Do not get confused with *two-phase locking*
- Two-phase commit can be used for managing distributed transactions in general
 - Example: A travel site that arranges both flight and hotel
 - * User wants to book the flight F and hotel H using credit card C
 - * The site needs to coordinate transaction with with banks, hotels, and airlines
 - * Q: How should the site handle the booking?

- * Q: What if the credit card authorizations fail? What if the flight is no longer available? What if the hotel is no longer available?

Asynchronous transaction

- Q: What if one participant is very slow?
 - Example: Starbucks. Cashier faster. Barista slower.
- Q: What does two-phase commit mean in this scenario?
- Q: How can we let each participant go ahead without waiting for the slow one?
- Q: What does Starbucks do?
- *Asynchronous transaction*
 - Each participant “commits” whenever he is done and moves ahead
 - * Transaction = sequence of smaller transactions by each participant
 - The entire transaction is done when every participant commits
 - No coordinated wait and synchronous commit
- Q: What if the coffee machine breaks down after customer paid?
 - *Compensating transaction*
 - * A transaction that “rolls back” a committed transaction
 - The coordinator should keep track of the “dependency” of transactions
 - * Together with their compensating transactions
 - If any transaction aborts, run compensating transaction for all committed transactions
- Q: When should we use two-phase commit/asynchronous transaction?
 - Importance of individual commit guarantee

- Duration of individual transaction
- Probability of abort