

CS144: Content Encoding

MIME (Multi-purpose Internet Mail Extensions)

- Q: Only “bits” are transmitted over the Internet. How does a browser/application interpret the bits and display them correctly?
- MIME (Multi-purpose Internet Mail Extensions)
 - Standard ways to “transmit” multimedia content over the Internet
 - Originally developed for email attachments, but currently used for all Internet data transmission
- MIME types
 - Specified as “type/subtype”. RFC2046 standard.
 - * IANA (Internet Assigned Number Authority) manages the official registry of all media types
 - In HTTP, it is specified in “Content-Type” header
 - * E.g., `Content-Type: text/html`
 - Popular types:
 - * Text: `text/plain`, `text/html`, `text/css`, ...
 - * Image: `image/jpeg`, `image/png`, `image/gif`, ...
 - * Audio: `audio/mpeg` (.mp3), `audio/mp4` (.mpa), ...
 - * Video: `video/mp4`, `video/H264`, `video/x-flv`, ...
 - * Application: `application/pdf`, `application/octet-stream`, ...
 - * Multipart: more on this later
- Q: What multimedia types/format should a browser support?
 - HTML5 is content-type/codec agnostic
 - No particular format support is required
 - Users expect browsers support “popular” codecs, such as JPG, PNG for images
 - Important legal issues regarding patent licensing
 - * 1999 UNISYS patent claim on GIF (expired in 2003)
 - * Licensing uncertainty for H.264 internet streaming
 - In 1997, MPEG/LA was formed to facilitate patent license after

MPEG2 standardization

- MPEG/LA's position on H.264 internet streaming: free through 2010!
- Google's purchase of VP8 patents (On2) in 2010
- * H.265 (High Efficiency Video Coding): MPEG/LA vs HEVC Advance
 - Google's push for AV1 (AOMedia Video 1: Alliance for Open Media Video 1)

Text Encoding

- Q: How does a browser map a sequence bits to characters if it is text?
- Character encoding/Character set
 - Mapping between numeric numbers and alphabetic characters
 - Many different character encodings
- Early character encodings (until mid 90's)
 - ASCII (American Standard Code for Information Interchange)
 - * created in 1963. First published as standard in 1967.
 - * 7bits. defines codes for 128 characters
 - * the basis of most of current encoding of roman characters
 - EBCDIC (Extended Binary Coded Decimal Interchange Code)
 - * created in 1963 by IBM for IBM mainframes
 - * 8bits. designed to be easy to represent in punch cards
 - * still used by some IBM mainframes.
 - ISO-8859-1 (= Latin-1)
 - * 8bits. consisting of 191 characters from the Latin script
 - * ASCII non-control characters have the same encoding
 - * used throughout Western Europe and America.
 - * ISO-8859-15, Windows-1252: more characters for French, Estonian,...
 - Local/regional encoding
 - * local character codes developed by each country
 - * DBCS (Double Byte Code Character Set)
 - one or two bytes are used to represent a character
 - frequently used in Asia

- Example: GB2312 (Simplified Chinese), EUC-KR (Korean), ...
- Q: What are the problems of multiple encoding standards?
- *Code page* (= character encoding)
 - a unique number given to a particular character encoding by a system
 - * On Windows: Hebrew (862), Greek (727), Korean (949)
 - OS sets the global code page for the computer
 - Q: What are the problems of a system-wide code-page setting?

UNICODE

- Motivation: Assign a unique number for every character in the world!
- International text encoding standard managed by Unicode Consortium
 - First standard, Unicode version 1.0 was published in October 1991
 - (almost) yearly release of a new Unicode version
- Every character maps to a CODE POINT
 - A -> U+0041
 - Hello -> U+0048 U+0065 U+006C U+006C U+006F.
- Originally defined to be a 16bit standard
 - No longer true. Currently 21bits (0x000000 – 0x10FFFFFF)
- a CODE POINT may be encoded into a sequence of bytes through an encoding scheme

UCS-2 (2-byte Universal Character Set)

- the first encoding scheme used for Unicode
- Represent the (original) unicode characters with two bytes
 - U+0041 -> 00 41
- Unicode byte order mark: U+FEFF
 - little endian/big endian issue
 - gives hints on the endian mode
 - stored at the beginning of a Unicode string

- Used by many systems, including Windows, macOS, Java, .NET, ...
- Q: Any problem with UCS-2 scheme?
 - space waste
 - Q: What will C program do for unicode-encoded data "a" (00 41)?
 - * Legacy applications cannot handle UNICODE data even if data contains only ASCII characters
 - Q: What will a UNICODE program do for the input 41 42 43 44?
 - * UNICODE applications cannot handle legacy ASCII data
- Q: If one byte is lost in the middle from 00 41 00 42 00 43, how is it interpreted?
- UCS-2 did not take off much for internet applications

UTF-8

- Primary goal: backward compatibility with ASCII encoding
 - Both UTF-8 and ASCII encoding should map all ASCII characters to the *same* 1-byte number
 - * e.g., A: U+0041 -> 41
 - * Q: What is the benefit of this property?
 - Allow easy recovery of the string from error
 - * even if a byte is missing, recover from the next character
- UTF-8 encoding standard
 - 0000 - 007F: 00000000 0zzzzzzz -> 0zzzzzzz
 - 0080 - 07FF: 00000yyy yyzzzzzz -> 110yyyyy 10zzzzzz
 - 0800 - FFFF: xxxxyyyy yyzzzzzz -> 1110xxxx 10yyyyyy 10zzzzzz
- Q: What will be UTF-8 encoding of character A (U+0041)?

- Example: 11010111 10111000 11101010 10111101 10110110 01111000
 - Q: How many characters in the example? How can we tell the beginning of a new character?
 - Q: How to recover if the second byte is lost during transmission?
- Q: If two strings are of the same length, are their encodings of the same length?
 - variable length encoding vs. fixed-length encoding
- UTF-8 encoding is the most popular encoding standard on Internet
 - Used by > 90% web sites

UTF-16

- Extension of UCS-2 after UNICODE became 21 bits
 - Two bytes are not enough to represent every unicode character!
 - Make the encoding as similar as possible to UCS-2, but allowing more characters
 - Variable length: either 2 bytes or 4 bytes
 - * U+0000 to U+D7FF and U+E000 to U+FFFF: 2-byte encoding just like UCS-2
 - * U+10000 to U+10FFFF: use 4-bytes to represent them
- Other Unicode encodings also exist
 - e.g., UTF-32: “32bit encoding”, ...

Using UNICODE

Q: How can we use/specify UNICODE?

- HTTP: Text type character encoding is specified as the “charset” parameter
 - E.g., Content-Type: text/html; charset=UTF-8
- HTML5:

- `<meta charset="utf-8">`
 - `)` for U+0041
 - For Web pages, UTF-8 encoding is by far the most popular encoding standard
- Most modern OS's support Unicode natively
 - Windows, macOS: UTF-16, Linux: UTF-8, ...
- Most modern languages, like Java and Javascript, use unicode as the default string type
 - provide multiple encoding/decoding functions for UTF-8, UTF-16, ISO-8859-1, ...
- But, unicode support in C++ is *messy*
 - On Linux, standard libraries, like `std::string`, support UTF-8
 - On Windows, UTF-16 can be used:
 - * use character type `wchar_t` (wide char) instead of `char`
 - * use wcs functions instead of str functions. (e.g., `wcslen` instead of `strlen`).
 - * prefix string constant with L, like L"Hello"

References

- MIME: RFC 2046
 - IANA media type list: <https://www.iana.org/assignments/media-types/media-types.xhtml>
- Unicode standard: <http://www.unicode.org/versions/latest/>